

Step 10: Analysis Questions

1. What is the MFLOP/s performance gain going from the CPU-only code to the final version of your CUDA code (the one with the `cudaMemPrefetchAsync()` call)? Please report this gain in terms of a multiplier, e.g., 2.5x, rather than with an absolute number or a percentage. Show your work on how you compute this result.

Comparing the CPU only MFLOPS to the final version of cuda we would do

CPU only = 553.66
Final version = 462.751

$$462.751 / 553.66 = 0.8358$$

So that means that the MFLOP/s performance gain would be 0.8358x

2. What is the memory bandwidth performance gain (or loss) going from the CPU-only code to the final version of your CUDA code (the one with the `cudaMemPrefetchAsync()` call)? Please report this gain in terms of a multiplier, e.g., 2.5x, rather than with an absolute number or a percentage. Show your work on how you compute this result.

$$5.83 / 6.96 = 0.837$$

There is performance loss of 0.837x

3. For the final version of your CUDA code (the one with the `cudaMemPrefetchAsync()` call), what is the total number of concurrent threads being run? Show your work on how you arrive at this result.

So we declared that the block size is 256, there a kernel with 1 block because of
`add<<<1, 256>>>(N, x, y);`

So it would be number of blocks * threads in a block so it would be $1 * 256 = 256$. Therefore the final cuda code has 256 concurrent threads

Table

A	B	C	D
tion	Elapsed time (ms)	MFLOP/s	memory bandwidth
Addition	882.14	553.66	6.96
or Addition: 1 thread, 1 thread block	51,829.42	9.4209	0.118
or Addition: 256 threads, 1 thread block	1,684.88	289.801	3.65
or Addition: 256 threads/block, many blocks	1,675.01	291.509	3.67
or Addition: 256 threads/block, many blocks, explicit data movement	1,055.17	462.751	5.83
		(ms)/1000 = seconds	
		(# arithmetic ops/ seconds) / 1024*1024	